

▼ Data Preprocessing

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('/content/drive/MyDrive/dataset')
df.head()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	Postal Code	Lattitud
0	6762810145	42491	5	2.50	3650	9050	2.0	0	4	5	...	1921	0	122003	52.864
1	6762810635	42491	4	2.50	2920	4000	1.5	0	0	5	...	1909	0	122004	52.887
2	6762810998	42491	5	2.75	2910	9480	1.5	0	0	3	...	1939	0	122004	52.885
3	6762812605	42491	4	2.50	3310	42998	2.0	0	0	3	...	2001	0	122005	52.953
4	6762812919	42491	3	2.00	2710	4500	1.5	0	0	4	...	1929	0	122006	52.904

Saved successfully!



```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
df.tail()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	Postal Code	Latt
14615	6762830250	42734	2	1.5	1556	20000	1.0	0	0	4	...	1957	0	122066	52
14616	6762830339	42734	3	2.0	1680	7000	1.5	0	0	4	...	1968	0	122072	52
14617	6762830618	42734	2	1.0	1070	6120	1.0	0	0	3	...	1962	0	122056	52
14618	6762830709	42734	4	1.0	1030	6621	1.0	0	0	4	...	1955	0	122042	52
14619	6762831463	42734	3	1.0	900	4770	1.0	0	0	3	...	1969	2009	122018	52

5 rows × 23 columns



```
df.shape
```

(14620, 23)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               14620 non-null   int64  
 1   Date              14620 non-null   int64  
 2   number of bedrooms 14620 non-null   int64  
 3   number of bathrooms 14620 non-null   float64 
 4   living area        14620 non-null   int64  
 5   lot area            14620 non-null   int64
```

```

6  number of floors          14620 non-null float64
7  waterfront present       14620 non-null int64
8  number of views          14620 non-null int64
9  condition of the house   14620 non-null int64
10 grade of the house        14620 non-null int64
11 Area of the house(excluding basement) 14620 non-null int64
12 Area of the basement     14620 non-null int64
13 Built Year               14620 non-null int64
14 Renovation Year          14620 non-null int64
15 Postal Code              14620 non-null int64
16 Latitude                 14620 non-null float64
17 Longitude                14620 non-null float64
18 living_area_renov        14620 non-null int64
19 lot_area_renov           14620 non-null int64
20 Number of schools nearby 14620 non-null int64
21 Distance from the airport 14620 non-null int64
22 Price                    14620 non-null int64
dtypes: float64(4), int64(19)
memory usage: 2.6 MB

```

```
df.isnull().any()
```

id	False
Date	False
number of bedrooms	False
number of bathrooms	False
living area	False
lot area	False
condition of the house	False
grade of the house	False
Area of the house(excluding basement)	False
Area of the basement	False
Built Year	False
Renovation Year	False
Postal Code	False
Latitude	False
Longitude	False
living_area_renov	False
lot_area_renov	False
Number of schools nearby	False
Distance from the airport	False
Price	False
dtype: bool	

Missing values

```
# if we have null values in the categorical columns we have to replace it with mode
```

```

## For the numerical column -null values replacement
## df['lotarea'].fillna['lotarea'].median(),inplace=True
##or
## df['lotarea'].fillna['lotarea'].mean(),inplace=True()

## For the categorical column-null values replacement

## df['lotarea'].fillna(df['lotarea'].mode(),inplace=True)

```

Descriptive statistic

```
df.head()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	Postal Code	Lattitud
0	6762810145	42491	5	2.50	3650	9050	2.0	0	4	5	...	1921	0	122003	52.864
1	6762810635	42491	4	2.50	2920	4000	1.5	0	0	5	...	1909	0	122004	52.887

```
df.Date.unique()
```

```
array([42491, 42492, 42493, 42494, 42495, 42496, 42497, 42498, 42499,
       42500, 42501, 42502, 42503, 42504, 42505, 42506, 42507, 42508,
       42509, 42510, 42511, 42512, 42513, 42514, 42515, 42516, 42517,
       42518, 42519, 42520, 42521, 42522, 42523, 42524, 42525, 42526,
       42527, 42528, 42529, 42530, 42531, 42532, 42533, 42534, 42535,
       42536, 42537, 42538, 42539, 42540, 42541, 42542, 42543, 42544,
       42545, 42546, 42547, 42548, 42549, 42550, 42551, 42552, 42553,
       42554, 42555, 42556, 42557, 42558, 42559, 42560, 42561, 42562,
       42563, 42564, 42565, 42566, 42567, 42568, 42569, 42570, 42571,
       42572, 42573, 42574, 42575, 42576, 42577, 42578, 42579, 42580,
       42581, 42582, 42583, 42584, 42585, 42586, 42587, 42588, 42589,
       42590, 42591, 42592, 42593, 42594, 42595, 42596, 42597, 42598,
       42599, 42600, 42601, 42602, 42603, 42604, 42605, 42606, 42607,
       42608, 42609, 42610, 42611, 42612, 42613, 42614, 42615, 42616,
       42617, 42618, 42619, 42620, 42621, 42622, 42623, 42624, 42625,
       42626, 42627, 42628, 42629, 42630, 42631, 42632, 42633, 42634,
       42635, 42636, 42637, 42638, 42639, 42640, 42641, 42642, 42643,
       42644, 42645, 42646, 42647, 42648, 42649, 42650, 42651, 42652,
       42653, 42654, 42655, 42656, 42657, 42658, 42659, 42660, 42661,
       42662, 42663, 42664, 42665, 42666, 42667, 42668, 42669, 42670,
       42671, 42672, 42673, 42674, 42675, 42676, 42677, 42678, 42679,
       42680, 42681, 42682, 42683, 42684, 42685, 42686, 42687, 42688,
       42689, 42690, 42691, 42692, 42693, 42694, 42695, 42696, 42697,
       42698, 42699, 42701, 42702, 42703, 42704, 42705, 42706, 42707,
       42708, 42709, 42710, 42711, 42712, 42713, 42714, 42715, 42716,
       42717, 42718, 42719, 42720, 42721, 42722, 42723, 42724, 42725,
       42726, 42727, 42728, 42729, 42730, 42731, 42732, 42733, 42734])
```

Saved successfully!

```
df.id.value_counts()
```

```
6762810145    1
6762812541    1
6762813736    1
6762813785    1
6762814009    1
...
6762812356    1
6762812410    1
6762812471    1
6762812563    1
6762831463    1
Name: id, Length: 14620, dtype: int64
```

```
df.Longitude.value_counts()
```

```
-114.552    73
-114.562    73
-114.490    72
-114.480    69
-114.553    68
...
-114.036    1
-114.096    1
-113.994    1
-114.111    1
-113.891    1
Name: Longitude, Length: 716, dtype: int64
```

```
df.Latitude.value_counts()
```

```
52.7745    13
52.8947    13
52.9255    13
52.9146    13
52.8834    13
...
52.8969    1
52.6834    1
52.5337    1
52.6155    1
```

```
52.5338      1
Name: Latitude, Length: 4662, dtype: int64
```

```
df.corr()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year
id	1.000000	0.045966	-0.329034	-0.516909	-0.648127	-0.100269	-0.312305	-0.112937	-0.293004	-0.045061	...	-0.068645
Date	0.045966	1.000000	-0.015663	-0.026485	-0.021958	0.004392	-0.010335	0.012006	-0.004782	-0.027402	...	-0.005869
number of bedrooms	-0.329034	-0.015663	1.000000	0.509784	0.570526	0.034416	0.177294	-0.006257	0.078665	0.026597	...	0.152954
number of bathrooms	-0.516909	-0.026485	0.509784	1.000000	0.753517	0.080806	0.502924	0.060104	0.183789	-0.128232	...	0.498127
living area	-0.648127	-0.021958	0.570526	0.753517	1.000000	0.174420	0.354743	0.105837	0.287728	-0.063358	...	0.309602
lot area	-0.100269	0.004392	0.034416	0.080806	0.174420	1.000000	-0.004138	0.026282	0.078308	-0.008548	...	0.051615
number of floors	-0.312305	-0.010335	0.177294	0.502924	0.354743	-0.004138	1.000000	0.016316	0.020153	-0.269928	...	0.481565
waterfront present	-0.112937	0.012006	-0.006257	0.060104	0.105837	0.026282	0.016316	1.000000	0.400206	0.018644	...	-0.024226
number of views	-0.293004	0.004782	0.078665	0.183789	0.287728	0.078308	0.020153	0.400206	1.000000	0.052533	...	-0.055357
Saved successfully!												
house												
grade of the house	-0.673448	-0.033097	0.352945	0.663054	0.761835	0.110546	0.463082	0.079831	0.254532	-0.152530	...	0.440358
Area of the house(excluding basement)	-0.565116	-0.015994	0.473599	0.684391	0.875793	0.183553	0.525643	0.071865	0.162672	-0.167695	...	0.419369
Area of the basement	-0.290806	-0.015711	0.300332	0.287190	0.441491	0.019755	-0.242976	0.085441	0.293062	0.180609	...	-0.138843
Built Year	-0.068645	-0.005869	0.152954	0.498127	0.309602	0.051615	0.481565	-0.024226	-0.055357	-0.381718	...	1.000000
Renovation Year	-0.109155	-0.011636	0.016132	0.049669	0.059400	0.006848	0.006705	0.085865	0.102944	-0.062126	...	-0.233683
Postal Code	0.294709	0.018243	-0.044156	-0.105546	-0.080303	0.070131	-0.129788	0.038318	0.039268	0.045334	...	-0.062349
Latitude	-0.479334	-0.023327	-0.013163	0.031156	0.054518	-0.090983	0.050731	-0.021795	-0.004555	-0.002998	...	-0.143153
Longitude	-0.070841	-0.018231	0.135712	0.223904	0.240208	0.221432	0.127550	-0.047791	-0.079706	-0.121189	...	0.414591
living_area_renov	-0.599900	-0.032495	0.389855	0.570530	0.757571	0.149744	0.285093	0.085743	0.281452	-0.099743	...	0.328625
lot_area_renov	-0.089604	-0.000050	0.029400	0.078627	0.180312	0.706812	-0.010120	0.032055	0.072300	-0.004748	...	0.072874
Number of schools nearby	-0.004821	-0.004071	0.003397	0.002180	0.002370	-0.012671	-0.007579	0.001563	0.008004	-0.006939	...	-0.001631
Distance from the airport	-0.004542	0.011457	-0.006157	0.009206	0.002511	0.003291	0.016567	0.001448	-0.001657	-0.002136	...	-0.003968
Price	-0.773114	-0.027919	0.308460	0.531735	0.712169	0.081992	0.262732	0.263687	0.395973	0.041376	...	0.050307

23 rows × 23 columns

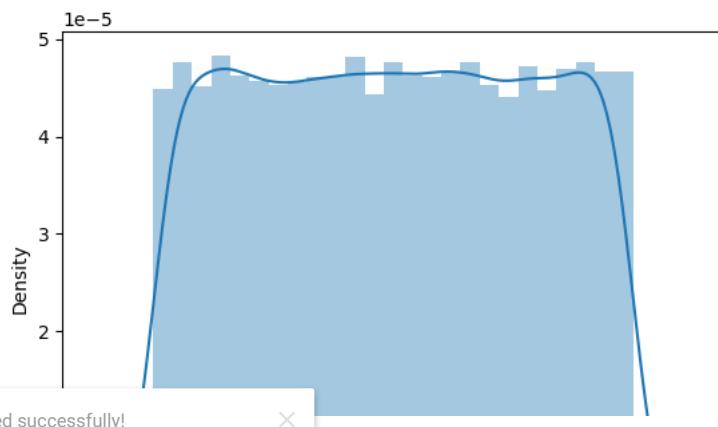


Univariate analysis

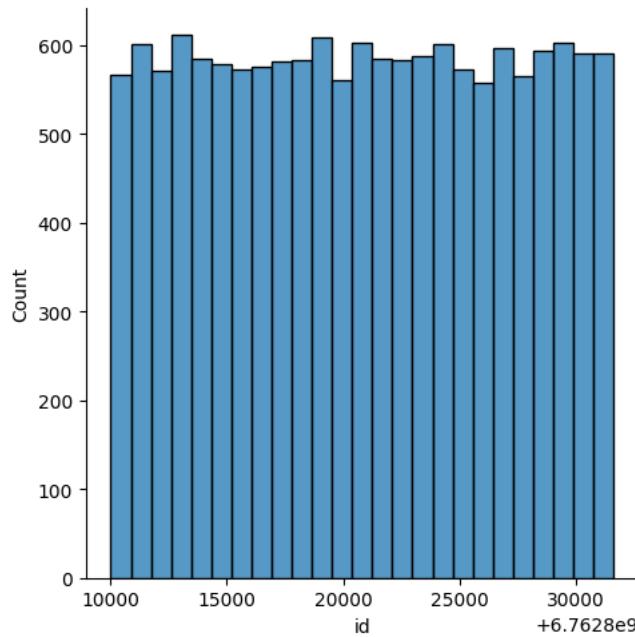
```
sns.distplot(df.id)
```

```
<ipython-input-25-ae12bbc1c629>:1: UserWarning:  
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.  
Please adapt your code to use either `displot` (a figure-level function with  
similar flexibility) or `histplot` (an axes-level function for histograms).  
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

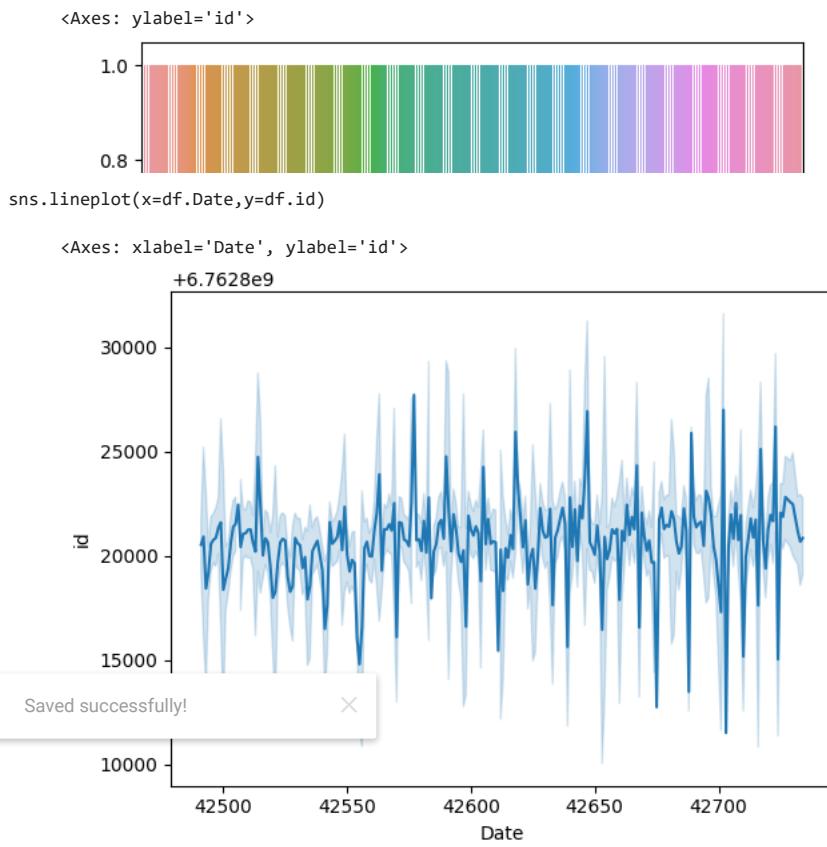
```
sns.distplot(df.id)  
<Axes: xlabel='id', ylabel='Density'>
```



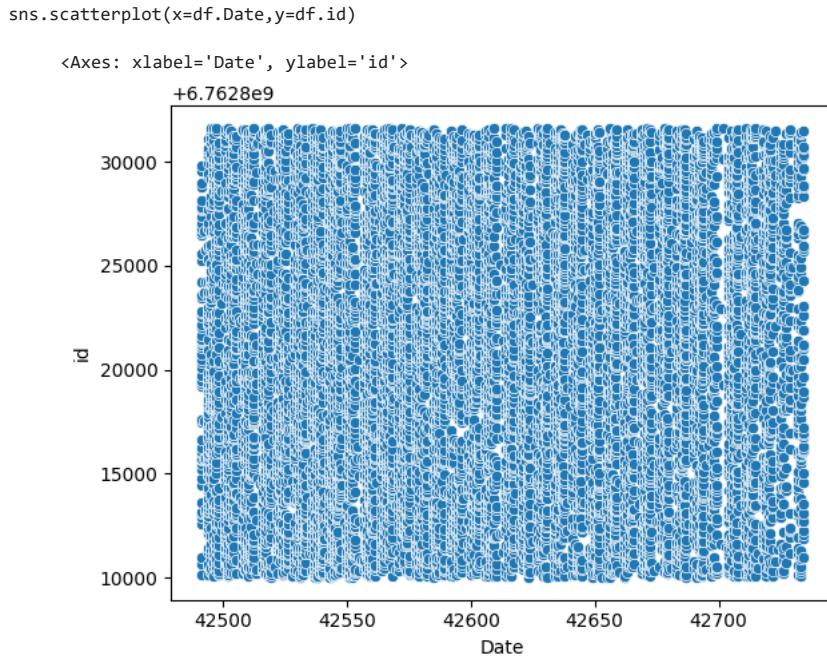
```
<seaborn.axisgrid.FacetGrid at 0x7fa3bc2ff0a0>
```



```
sns.barplot(x=df.id.value_counts().index,y=df.id.value_counts())
```



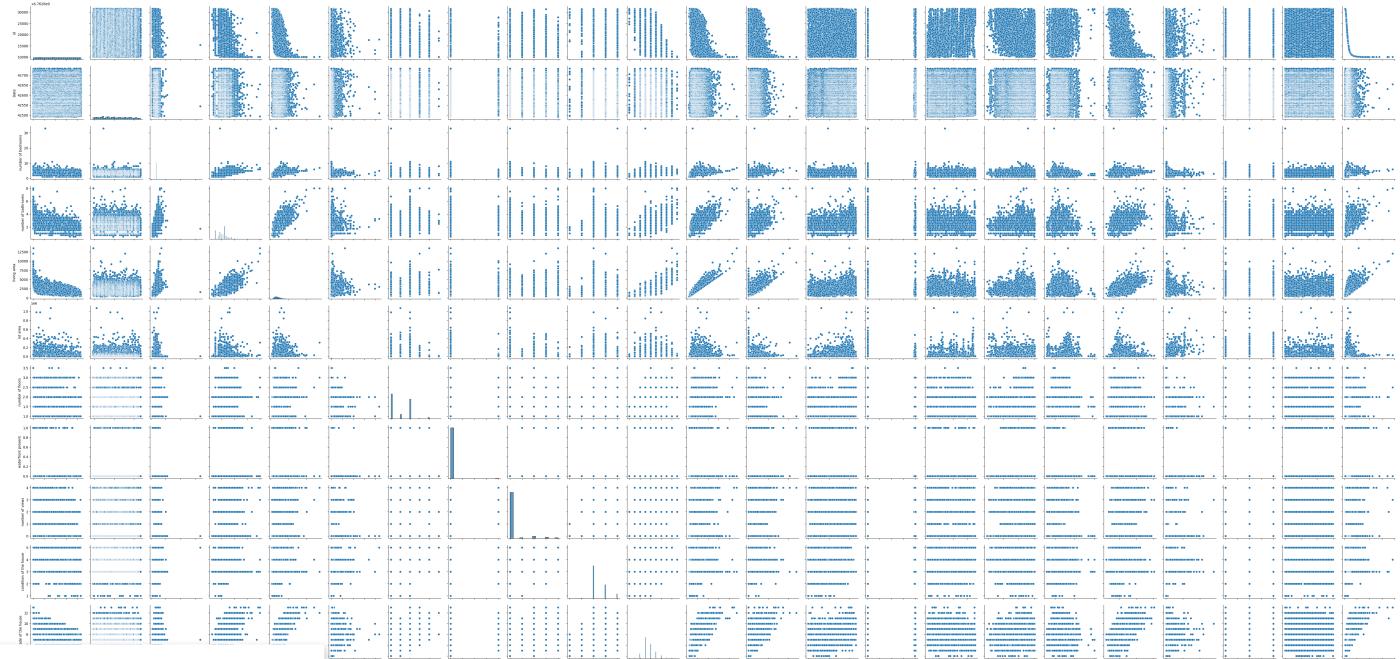
▼ Bivariate Analysis



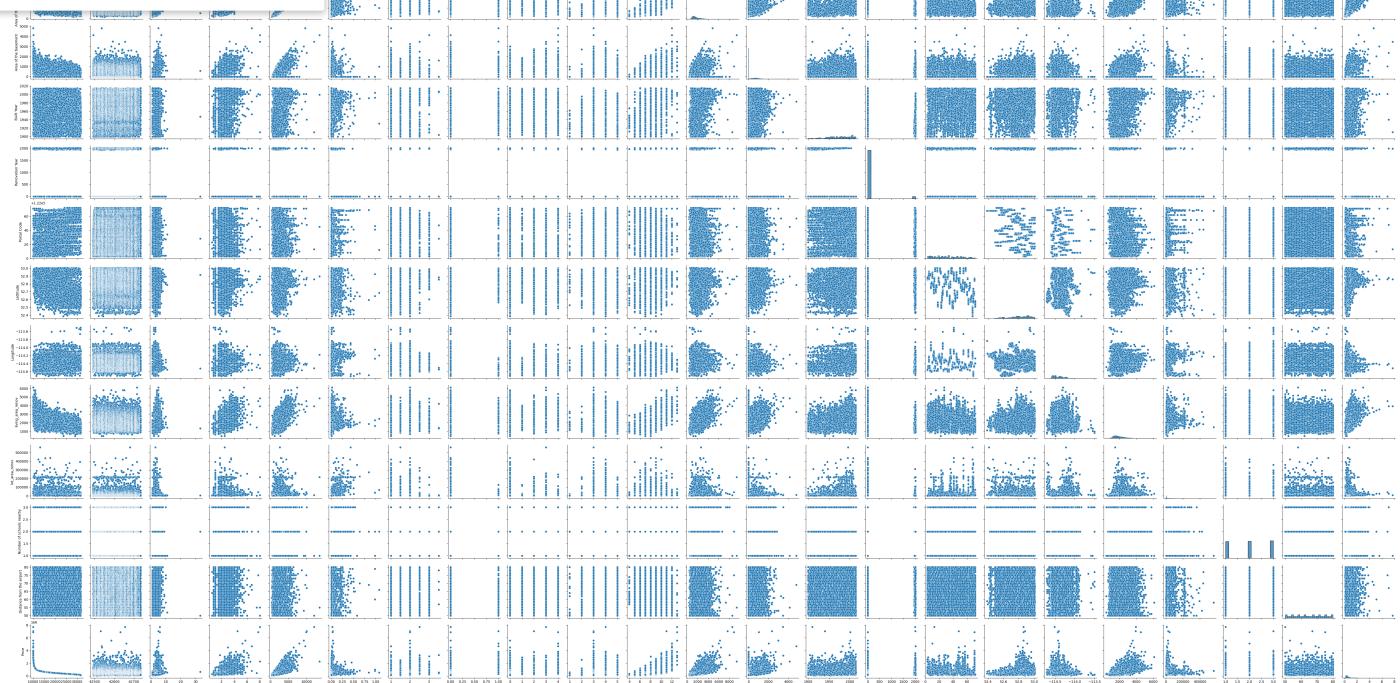
Multivariate Analysis

```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7fa3b6a1a830>
```

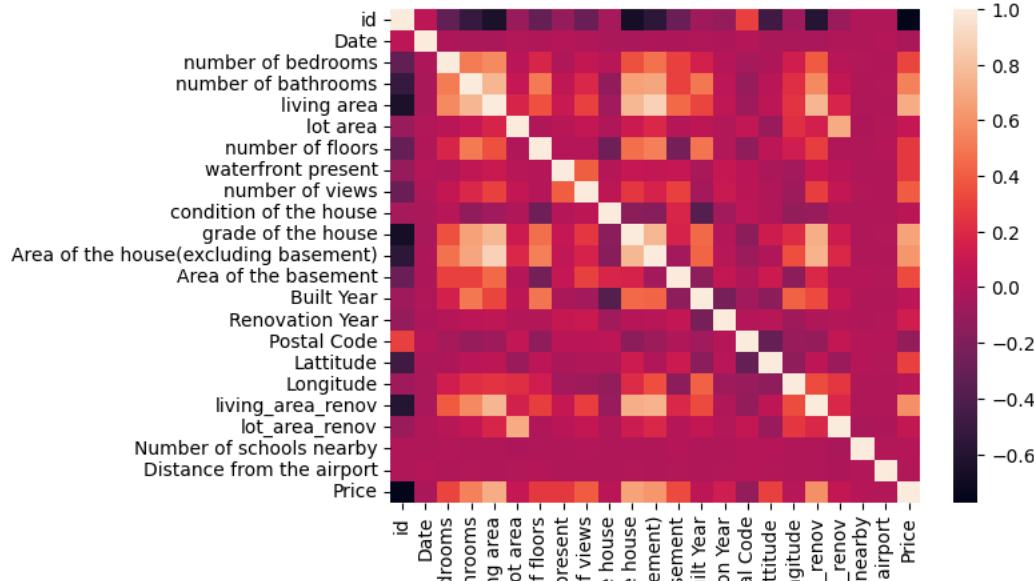


Saved successfully!



```
sns.heatmap(df.corr())
```

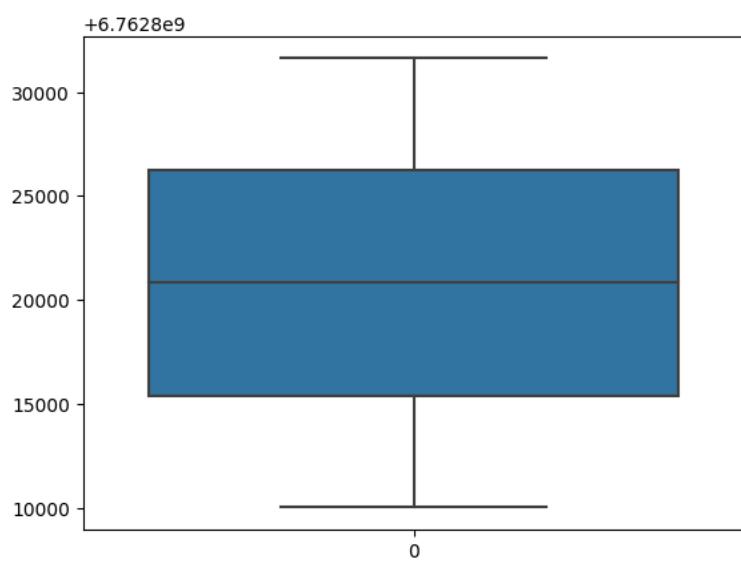
<Axes: >



Outlier -Detection

Saved successfully!

<Axes: >

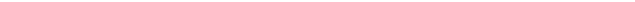


▼ Bivariate Analysis

```
sns.lineplot(x=df.Date,y=df.Latitude)
```

The figure is a line plot titled "Latitude vs Date". The y-axis is labeled "Latitude" and ranges from 52.85 to 53.05 with major ticks at 52.9 and 53.0. The x-axis is labeled "Date" and shows dates from January 1, 2010, to December 31, 2010. The data is represented by a blue line that exhibits high-frequency oscillations. The mean value of the latitude over time is approximately 52.95.

```
## Outlier removal -IQR method,Z-score,Percentile
```

—  —

Outlier-Replacement-Median

```
q1 = df.id.quantile(0.25) #(Q1)  
q3 = df.id.quantile(0.75) #(Q3)
```

52.5 1

```
print(q1)  
print(q3)
```

6762815404.75
6762826251.25

$$\text{IQR} = q_3 - q_1$$

Saved successfully!

lower limit = q1 - 1.5 * IQR

Upper_limit

6762842521.0

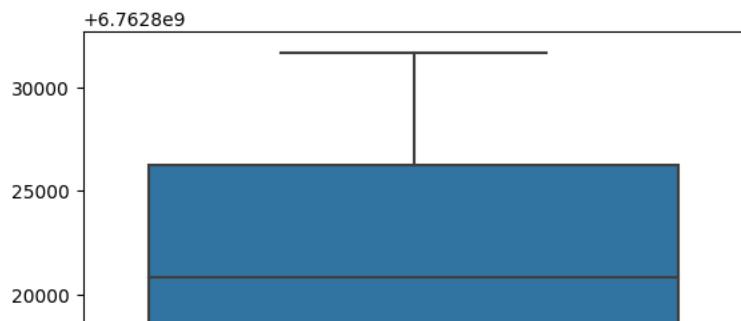
`df.median()`

```
id          6.762821e+09
Date        4.260000e+04
number of bedrooms 3.000000e+00
number of bathrooms 2.250000e+00
living area   1.930000e+03
lot area      7.620000e+03
number of floors 1.500000e+00
waterfront present 0.000000e+00
number of views 0.000000e+00
condition of the house 3.000000e+00
grade of the house 7.000000e+00
Area of the house(excluding basement) 1.580000e+03
Area of the basement 0.000000e+00
Built Year    1.975000e+03
Renovation Year 0.000000e+00
Postal Code   1.220320e+05
Latitude      5.280640e+01
Longitude     -1.144210e+02
living_area_renov 1.850000e+03
lot_area_renov 7.620000e+03
Number of schools nearby 2.000000e+00
Distance from the airport 6.500000e+01
Price         4.500000e+05
dtype: float64
```

```
df['id'] = np.where(df['id']>Upper_limit,30,df['id'])
```

```
sns.boxplot(df.id)
```

<Axes: >



df.shape

(14620, 23)

Encoding Technique

df.head()

	number	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	Postal Code	Lattitud
0	6762810145	42491	5	2.50	3650	9050	2.0	0	4	5	...	1921	0 122003 52.864
1	6762810635	42491	4	2.50	2920	4000	1.5	0	0	5	...	1909	0 122004 52.887
2	6762810998	42491	5	2.75	2910	9480	1.5	0	0	3	...	1939	0 122004 52.885
3	6762812605	42491	4	2.50	3310	42998	2.0	0	0	3	...	2001	0 122005 52.953
4	6762812919	42491	3	2.00	2710	4500	1.5	0	0	4	...	1929	0 122006 52.904

5 rows × 23 columns



label encoding

from sklearn.preprocessing import LabelEncoder

le=LabelEncoder()

df.id = le.fit_transform(df.id)
df.Date= le.fit_transform(df.Date)y=df.Date
y

0	0
1	0
2	0
3	0
4	0
...	
14615	240
14616	240
14617	240
14618	240
14619	240

Name: Date, Length: 14620, dtype: int64

df.head()

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	Postal Code	Latitude	Longitude
0	84	0	5	2.50	3650	9050	2.0	0	4	5	...	1921	0	122003	52.8645	-11
1	398	0	4	2.50	2920	4000	1.5	0	0	5	...	1909	0	122004	52.8878	-11
2	647	0	5	2.75	2910	9480	1.5	0	0	3	...	1939	0	122004	52.8852	-11
3	1735	0	4	2.50	3310	42998	2.0	0	0	3	...	2001	0	122005	52.9532	-11
4	1951	0	3	2.00	2710	4500	1.5	0	0	4	...	1929	0	122006	52.9047	-11

5 rows × 23 columns

▼ One hot Encoding Method

```
df1= pd.get_dummies(df,columns=['id'])
df1
```

	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	grade of the house	...	id_14610	id_14611	id_14612	id_14613
Saved successfully!														
0	0	5	2.50	3650	9050	2.0	0	4	5	10	...	0	0	0
1	0	4	2.50	2920	4000	1.5	0	0	5	8	...	0	0	0
2	0	5	2.75	2910	9480	1.5	0	0	3	8	...	0	0	0
3	0	4	2.50	3310	42998	2.0	0	0	3	9	...	0	0	0
4	0	3	2.00	2710	4500	1.5	0	0	4	8	...	0	0	0
...
14615	240	2	1.50	1556	20000	1.0	0	0	4	7	...	0	0	0
14616	240	3	2.00	1680	7000	1.5	0	0	4	7	...	0	0	0
14617	240	2	1.00	1070	6120	1.0	0	0	3	6	...	0	0	0
14618	240	4	1.00	1030	6621	1.0	0	0	4	6	...	0	0	0
14619	240	3	1.00	900	4770	1.0	0	0	3	6	...	0	0	0

14620 rows × 14642 columns



▼ X and y split

```
X=df.drop(columns=['Price'],axis = 1)
X.head()
```

```

      id Date number of bedrooms number of bathrooms living area lot area number of floors waterfront present number of views condition of the house ...
      0   2380000
      1   1400000
      2   1200000
      3   838000
      4   805000
      ...
      14615 221700
      14616 219200
      14617 209000
      14618 205000
      14619 146000
Name: Price, Length: 14620, dtype: int64

```

Scaling

```

## MinMax Scaling
Saved successfully! ×
scale =MinMaxScaler()
scale =MinMaxScaler()

X_scaled= pd.DataFrame(scale.fit_transform(X),columns = X.columns)

X_scaled.head()

```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Area of the basement	Built Year	Renovation Year	Postal Code	Latitude	Longitude
0	0.005746	0.0	0.12500	0.266667	0.249051	0.007945	0.4	0.0	1.0	1.00	...	0.058091	0.182609	0.0	0.0	52.3562	-71.0589
1	0.027225	0.0	0.09375	0.266667	0.193622	0.003241	0.2	0.0	0.0	1.00	...	0.209544	0.078261	0.0	0.0	52.3562	-71.0589
2	0.044257	0.0	0.12500	0.300000	0.192863	0.008345	0.2	0.0	0.0	0.50	...	0.000000	0.339130	0.0	0.0	52.3562	-71.0589
3	0.118681	0.0	0.09375	0.266667	0.223235	0.039562	0.4	0.0	0.0	0.50	...	0.000000	0.878261	0.0	0.0	52.3562	-71.0589
4	0.133456	0.0	0.06250	0.200000	0.177677	0.003707	0.2	0.0	0.0	0.75	...	0.172199	0.252174	0.0	0.0	52.3562	-71.0589

5 rows × 22 columns

Train and Split

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(X_scaled,y,test_size=0.3,random_state = 0)

x_train.shape
(10234, 22)

x_test.shape
(4386, 22)

y_test.shape

```

(4386,)

```
from keras.models import Sequential
from keras.layers import Dense

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

import matplotlib.pyplot as plt

import numpy as np

np.histogram([1, 2])

(array([1, 0, 0, 0, 0, 0, 0, 0, 1]),
 array([1., 1.1, 1.2, 1.3, 1.4, 1.5, 1.6, 1.7, 1.8, 1.9, 2. ]))

x = df.drop('Price', axis=1)
x.head()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Area of the basement	Built Year	Renovation Year	Postal Code	Lati...
Saved successfully!																
1	398	0	4	2.50	3650	9050	2.0	0	4	5	...	280	1921	0	122003	52.1
2	647	0	5	2.75	2910	9480	1.5	0	0	3	...	1010	1909	0	122004	52.1
3	1735	0	4	2.50	3310	42998	2.0	0	0	3	...	0	2001	0	122005	52.1
4	1951	0	3	2.00	2710	4500	1.5	0	0	4	...	830	1929	0	122006	52.1

5 rows × 22 columns



```
y = df['Price']
y

0      2380000
1      1400000
2      1200000
3      838000
4      805000
...
14615    221700
14616    219200
14617    209000
14618    205000
14619    146000
Name: Price, Length: 14620, dtype: int64
```

```
from sklearn.model_selection import train_test_split
```

```
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.2,random_state=12)
```

▼ Build an ANN Model

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14619 entries, 0 to 14619
Data columns (total 23 columns):
 #   Column          Non-Null Count  Dtype  

```

```
--- -----
0 id                           14620 non-null int64
1 Date                          14620 non-null int64
2 number of bedrooms           14620 non-null int64
3 number of bathrooms          14620 non-null float64
4 living area                   14620 non-null int64
5 lot area                      14620 non-null int64
6 number of floors              14620 non-null float64
7 waterfront present           14620 non-null int64
8 number of views               14620 non-null int64
9 condition of the house       14620 non-null int64
10 grade of the house          14620 non-null int64
11 Area of the house(excluding basement) 14620 non-null int64
12 Area of the basement        14620 non-null int64
13 Built Year                  14620 non-null int64
14 Renovation Year             14620 non-null int64
15 Postal Code                 14620 non-null int64
16 Latitude                     14620 non-null float64
17 Longitude                    14620 non-null float64
18 living_area_renov           14620 non-null int64
19 lot_area_renov              14620 non-null int64
20 Number of schools nearby    14620 non-null int64
21 Distance from the airport   14620 non-null int64
22 Price                        14620 non-null int64
dtypes: float64(4), int64(19)
memory usage: 2.6 MB
```

Saved successfully! X

```
Date                           0
number of bedrooms            0
number of bathrooms           0
living area                   0
lot area                      0
number of floors              0
waterfront present           0
number of views               0
condition of the house       0
grade of the house           0
Area of the house(excluding basement) 0
Area of the basement         0
Built Year                    0
Renovation Year               0
Postal Code                   0
Latitude                      0
Longitude                     0
living_area_renov             0
lot_area_renov                0
Number of schools nearby     0
Distance from the airport    0
Price                         0
dtype: int64
```

```
df['id'].unique()

array([ 84,  398,  647, ..., 13943, 14015, 14519])
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()

df['Date'].unique()

array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9,  10,  11,  12,
       13,  14,  15,  16,  17,  18,  19,  20,  21,  22,  23,  24,  25,
       26,  27,  28,  29,  30,  31,  32,  33,  34,  35,  36,  37,  38,
       39,  40,  41,  42,  43,  44,  45,  46,  47,  48,  49,  50,  51,
       52,  53,  54,  55,  56,  57,  58,  59,  60,  61,  62,  63,  64,
       65,  66,  67,  68,  69,  70,  71,  72,  73,  74,  75,  76,  77,
       78,  79,  80,  81,  82,  83,  84,  85,  86,  87,  88,  89,  90,
       91,  92,  93,  94,  95,  96,  97,  98,  99,  100, 101, 102, 103,
       104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116,
       117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129,
```

```
130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140, 141, 142,
143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154, 155,
156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168,
169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181,
182, 183, 184, 185, 186, 187, 188, 189, 190, 191, 192, 193, 194,
195, 196, 197, 198, 199, 200, 201, 202, 203, 204, 205, 206, 207,
208, 209, 210, 211, 212, 213, 214, 215, 216, 217, 218, 219, 220,
221, 222, 223, 224, 225, 226, 227, 228, 229, 230, 231, 232, 233,
234, 235, 236, 237, 238, 239, 240])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14620 entries, 0 to 14619
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               14620 non-null   int64  
 1   Date              14620 non-null   int64  
 2   number of bedrooms 14620 non-null   int64  
 3   number of bathrooms 14620 non-null   float64 
 4   living area       14620 non-null   int64  
 5   lot area          14620 non-null   int64  
 6   number of floors  14620 non-null   float64 
 7   waterfront present 14620 non-null   int64  
 8   number of views   14620 non-null   int64  
 9   condition of the house 14620 non-null   int64  
 10  grade of the house 14620 non-null   int64  
 11  area of the basement 14620 non-null   int64  
 12  renovation year  14620 non-null   int64  
 13  postal code       14620 non-null   int64  
 14  latitude          14620 non-null   float64 
 15  longitude          14620 non-null   float64 
 16  living_area_renov 14620 non-null   int64  
 17  lot_area_renov    14620 non-null   int64  
 18  number of schools nearby 14620 non-null   int64  
 19  distance from the airport 14620 non-null   int64  
 20  price              14620 non-null   float64(4) 
 21  memory usage       2.6 MB
```

Saved successfully!

```
x = df.drop('Price', axis=1)
x.head()
```

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Area of the basement	Built Year	Renovation Year	Postal Code	Lati
0	84	0	5	2.50	3650	9050	2.0	0	4	5	...	280	1921	0	122003	52.8
1	398	0	4	2.50	2920	4000	1.5	0	0	5	...	1010	1909	0	122004	52.8
2	647	0	5	2.75	2910	9480	1.5	0	0	3	...	0	1939	0	122004	52.8
3	1735	0	4	2.50	3310	42998	2.0	0	0	3	...	0	2001	0	122005	52.9
4	1951	0	3	2.00	2710	4500	1.5	0	0	4	...	830	1929	0	122006	52.9

5 rows × 22 columns



```
y = df['Price']
y

0      2380000
1      1400000
2      1200000
3      838000
4      805000
...
14615    221700
14616    219200
14617    209000
14618    205000
```

```
14619      146000
Name: Price, Length: 14620, dtype: int64
```

```
from sklearn.model_selection import train_test_split

xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.2,random_state=12)

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

reg = Sequential()
reg.add(Dense(4,activation='relu'))
reg.add(Dense(12,activation='relu'))
reg.add(Dense(32,activation='relu'))
reg.add(Dense(1,activation='linear'))

reg.compile(optimizer='adam',loss='mse',metrics=['mse'])
```

Saved successfully!

```
reg.fit(xtrain,ytrain,batch_size=10,epochs=25,validation_data=(xtest,ytest))

Epoch 1/25
1170/1170 [=====] - 4s 3ms/step - loss: 148238516224.0000 - mse: 148238516224.0000 - val_loss: 154917945344.0000
Epoch 2/25
1170/1170 [=====] - 3s 2ms/step - loss: 66609639424.0000 - mse: 66609639424.0000 - val_loss: 67012669440.0000
Epoch 3/25
1170/1170 [=====] - 3s 2ms/step - loss: 32393545728.0000 - mse: 32393545728.0000 - val_loss: 51977404416.0000
Epoch 4/25
1170/1170 [=====] - 4s 3ms/step - loss: 24200749056.0000 - mse: 24200749056.0000 - val_loss: 37960912896.0000
Epoch 5/25
1170/1170 [=====] - 3s 2ms/step - loss: 19964098560.0000 - mse: 19964100608.0000 - val_loss: 32543971328.0000
Epoch 6/25
1170/1170 [=====] - 3s 2ms/step - loss: 17879166976.0000 - mse: 17879166976.0000 - val_loss: 28656603136.0000
Epoch 7/25
1170/1170 [=====] - 3s 3ms/step - loss: 16022565888.0000 - mse: 16022565888.0000 - val_loss: 25155653632.0000
Epoch 8/25
1170/1170 [=====] - 7s 6ms/step - loss: 14798630912.0000 - mse: 14798630912.0000 - val_loss: 23124426752.0000
Epoch 9/25
1170/1170 [=====] - 4s 4ms/step - loss: 13942464512.0000 - mse: 13942464512.0000 - val_loss: 21679075328.0000
Epoch 10/25
1170/1170 [=====] - 5s 4ms/step - loss: 13321062400.0000 - mse: 13321062400.0000 - val_loss: 20917897216.0000
Epoch 11/25
1170/1170 [=====] - 7s 6ms/step - loss: 12731298816.0000 - mse: 12731298816.0000 - val_loss: 19960096768.0000
Epoch 12/25
1170/1170 [=====] - 5s 4ms/step - loss: 12406204416.0000 - mse: 12406204416.0000 - val_loss: 18496010240.0000
Epoch 13/25
1170/1170 [=====] - 3s 3ms/step - loss: 11933257728.0000 - mse: 11933257728.0000 - val_loss: 17950363648.0000
Epoch 14/25
1170/1170 [=====] - 3s 3ms/step - loss: 11688812544.0000 - mse: 11688813568.0000 - val_loss: 17430949888.0000
Epoch 15/25
1170/1170 [=====] - 3s 2ms/step - loss: 11484267520.0000 - mse: 11484267520.0000 - val_loss: 17036865536.0000
Epoch 16/25
1170/1170 [=====] - 3s 2ms/step - loss: 11198681088.0000 - mse: 11198681088.0000 - val_loss: 16964756480.0000
Epoch 17/25
1170/1170 [=====] - 2s 2ms/step - loss: 11129534464.0000 - mse: 11129534464.0000 - val_loss: 16533778432.0000
Epoch 18/25
1170/1170 [=====] - 3s 2ms/step - loss: 10773383168.0000 - mse: 10773383168.0000 - val_loss: 16317324288.0000
Epoch 19/25
1170/1170 [=====] - 3s 3ms/step - loss: 10875949056.0000 - mse: 10875949056.0000 - val_loss: 16965533696.0000
Epoch 20/25
1170/1170 [=====] - 3s 3ms/step - loss: 10398380032.0000 - mse: 10398380032.0000 - val_loss: 16015109120.0000
Epoch 21/25
1170/1170 [=====] - 2s 2ms/step - loss: 10635054080.0000 - mse: 10635054080.0000 - val_loss: 15396625408.0000
Epoch 22/25
1170/1170 [=====] - 3s 2ms/step - loss: 10403016704.0000 - mse: 10403016704.0000 - val_loss: 16796754944.0000
Epoch 23/25
1170/1170 [=====] - 2s 2ms/step - loss: 10356554752.0000 - mse: 10356554752.0000 - val_loss: 17184329728.0000
Epoch 24/25
1170/1170 [=====] - 3s 3ms/step - loss: 10177088512.0000 - mse: 10177088512.0000 - val_loss: 16087503872.0000
Epoch 25/25
```

```
1170/1170 [=====] - 3s 3ms/step - loss: 10232511488.0000 - mse: 10232511488.0000 - val_loss: 16210734080.0000 - <keras.callbacks.History at 0x7fa30988a5f0>
```

```
reg.predict([[123123,345345,2343241,2]])  
  
1/1 [=====] - 1s 626ms/step  
array([[806956.7]], dtype=float32)  
  
# Handle missing values  
df.fillna(0,inplace=True)  
  
# Split the dataset into input and target variables  
X = df.iloc[:, :-1].values  
y = df.iloc[:, -1].values  
  
# Split the dataset into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)  
  
# Build the ANN model  
model = Sequential()  
Saved successfully!  
model.add(Dense(units=1, activation='relu'))  
model.add(Dense(units=1, activation='linear'))  
  
# Compile the model  
model.compile(optimizer='adam', loss='mean_squared_error')  
  
# Train the model  
model.fit(X_train, y_train, batch_size=32, epochs=100)
```

```
366/366 [=====] - 1s 2ms/step - loss: 108750046/2.0000
Epoch 93/100
366/366 [=====] - 1s 3ms/step - loss: 10796452864.0000
Epoch 94/100
366/366 [=====] - 1s 3ms/step - loss: 10686840832.0000
Epoch 95/100
366/366 [=====] - 1s 3ms/step - loss: 10557347840.0000
Epoch 96/100
366/366 [=====] - 1s 4ms/step - loss: 10259372032.0000
Epoch 98/100
366/366 [=====] - 1s 4ms/step - loss: 10175884288.0000
Epoch 99/100
366/366 [=====] - 1s 3ms/step - loss: 10178080768.0000
Epoch 100/100
366/366 [=====] - 1s 2ms/step - loss: 9940421632.0000
<keras.callbacks.History at 0x7fa309f51ed0>
```

▼ Test the model

```
import tensorflow as tf

# Test the model
y_pred = model.predict(X_test)
mse = tf.keras.losses.mean_squared_error(y_test, y_pred).numpy()
print('Mean Squared Error:', mse)

Saved successfully! [X]
```

=====] - 0s 1ms/step
+11 9.4334589e+11 1.5283644e+11 ... 1.3892839e+11 2.0528737e+11

✓ 0s completed at 1:58PM

